

IDS2Eval Scorecard

cic-ids2018

7 ok · 3 warning · 2 flag · 12 checks

REVIEW REQUIRED

- dedup_check

ok

label_conflict_check

ok

leakage_screen

ok

one_rule_check

ok

identity_column_flag

flag

temporal_leakage_check

ok

homogeneity_test

flag

class_distribution_report

warning

low_cardinality_warning

ok

schema_fingerprint_check

ok

data_integrity_check

warning

known_issue_lookup

warning
- train duplicates: 0 (0.00%); test rows leaking a train feature-match: 0 (0.00%); test-internal duplicates: 0

no feature vector maps to more than one label

no single feature or pair dominates importance (top1=8.6%, top2=15.4%)

single rule "TotLen Fwd Pkts <= 21" reaches 86.9% test accuracy (86.2% train)

standalone AUC by column: {'Dst Port': 0.912}. Suggest dropping: ['Dst Port']

no schema.timestamp_column configured

9 classes tested; leakage signature (test significantly closer than control, p<0.05) in: ['Bot', 'DoS attacks-Slowloris']

15 classes, train imbalance ratio (majority:minority) = 149712:1; classes below 1% of train: ['SSH-Bruteforce', 'DoS attacks-GoldenEye', 'DoS attacks-Slowloris', 'DDOS attack-LOIC-UDP', 'DoS attacks-SlowHTTPTest', 'FTP-BruteForce', 'Brute Force -Web', 'Brute Force -XSS', 'SQL Injection']


unique values by column: {'Dst Port': 24978}

no known extractor signature matched

missing values in 1 feature column(s); +-inf values in 2 feature column(s); 8 constant/near-constant feature(s)

The 2018-02-23 day-file's Brute-Force-Web/Brute-Force-XSS rows are approximately 41% mislabeled (ground-truth labeling error, not an extraction or pipeline bug). (Liu et al. 2022, IEEE CNS)

Class distribution (log scale)



Top feature importances (leakage_screen)

