

IDS2Eval Scorecard

unsw-nb15

16 ok · 2 warning · 1 flag · 19 checks

REVIEW REQUIRED

dedup_check	train duplicates: 0 (0.00%); test rows leaking a train feature-match: 0 (0.00%); test-internal duplicates: 0	ok
label_conflict_check	no feature vector maps to more than one label	ok
near_duplicate_class_check	no near-zero-distance feature vectors found across 10 classes tested	ok
leakage_screen	no single feature or pair dominates importance (top1=15.5%, top2=28.4%)	ok
one_rule_check	single rule 'sttl <= 61' reaches 68.0% test accuracy (88.1% train), vs. 63.9% from always predicting the majority class	ok
feature_auc_ranking_check	'sttl' alone reaches AUC 0.872 identifying '0.0'	ok
identity_column_flag	no id_like_columns configured	ok
port_protocol_shortcut_check	no port-like entry in schema.id_like_columns and/or no proto-like column found, this check needs both	ok
temporal_leakage_check	no schema.timestamp_column configured	ok
temporal_realism_check	no schema.timestamp_column configured	ok
flow_group_leakage_check	no schema.flow_id_columns configured	ok
row_order_leakage_check	adjacent-row label-transition rate (train=0.09, test=0.00) as a share of what a randomly shuffled ordering would produce; test is the most contiguous, consistent with rows still in collection order rather than shuffled before this split	flag
homogeneity_test	10 classes tested; inconclusive (95% CI straddles the +/-0.05 equivalence margin) for: ['Backdoor', 'Worms']; sample size isn't enough to tell for these classes	warning
scenario_holdout_falsification	no schema.scenario_column configured	ok
class_distribution_report	2 classes, train imbalance ratio (majority:minority) = 1:1	ok
low_cardinality_warning	no id_like_columns configured	ok
schema_fingerprint_check	no known extractor signature matched	ok
data_integrity_check	no missing labels, constant features, or +-inf values found	ok
known_issue_lookup	Some redistributions invert the UNSW_NB15_training-set.csv/testing-set.csv file-name-to-content mapping relative to the dataset's documented convention (82,332 vs. 175,341 rows), verify against the published per-attack-category counts before trusting file names. (Independently verified in this project's own audit methodology)	warning

Class distribution (log scale)



Top feature importances (leakage_screen)

